

# Course Review

Instructor: Matei Zaharia

[cs245.stanford.edu](https://cs245.stanford.edu)

# What Are Data-Intensive Systems?

**Relational databases:** most popular type of data-intensive system (MySQL, Oracle, etc)

**Many systems facing similar concerns:** message queues, key-value stores, streaming systems, ML frameworks, **your custom app?**

**Goal:** learn the main issues and principles that span all data-intensive systems

# Typical System Challenges

**Reliability** in the face of hardware crashes, bugs, bad user input, etc

**Concurrency:** access by multiple users

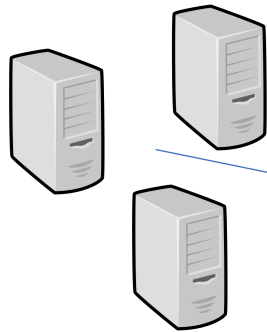
**Performance:** throughput, latency, etc

**Access interface** from many, changing apps

**Security** and data privacy

# Basic Components

Clients / users



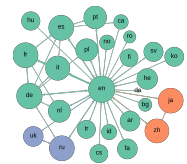
Queries



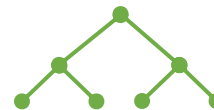
Administrator

Logical dataset  
(e.g. table, graph)

First Name	Last Name	Address	City	Age
Hickey	House	123 Fantasy Way	Anaheim	73
Bat	Man	321 Cavern Ave	Gotham	54
Wonder	Woman	987 Truth Way	Paradise	39
Donald	Duck	555 Quack Street	Mallard	65
Bugs	Bunny	567 Carrot Street	Rascal	58
Wiley	Coyote	999 Acme Way	Canyon	61
Cat	Woman	234 Purrfect Street	Hairball	32
Tweety	Bird	543	Ittobavi	28



Physical storage  
(data structures)



# Two Big Ideas

## Declarative interfaces

- » Apps specify *what* they want, not *how* to do it
- » Example: “store a table with 2 integer columns”, but not how to encode it on disk
- » Example: “count records where column1 = 5”

## Transactions

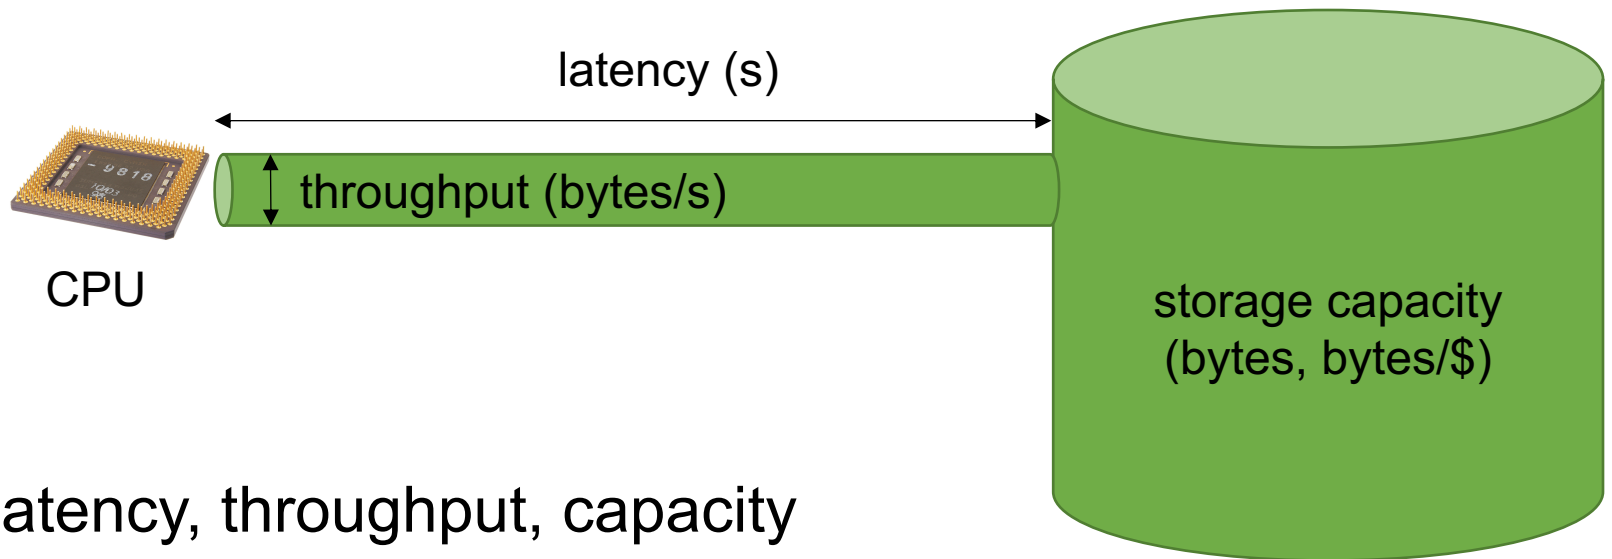
- » Encapsulate multiple app actions into one *atomic* request (fails or succeeds as a whole)
- » Concurrency models for multiple users
- » Clear interactions with failure recovery

# Key Concepts: Architecture

**Traditional RDBMS:** self-contained end to end system

**Data lake:** separate storage from compute engines to let many engines use same data

# Key Concepts: Hardware



Latency, throughput, capacity

Random vs sequential I/Os

Caching & 5-minute rule

# Key Concepts: Data Storage

**Field encoding**

**Record encoding:** fixed/variable format, etc

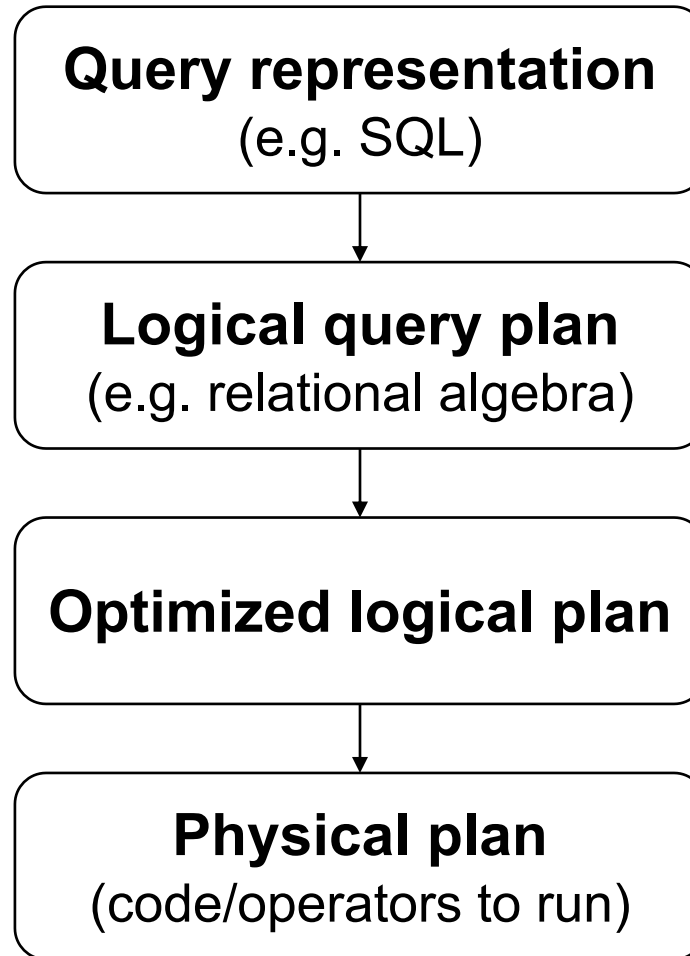
**Table encoding:** row or column oriented

**Data ordering**

**Indexes:** dense, sparse, B+ trees, hashing, multi-dimensional



# Key Concepts: Query Execution



Many **execution methods**: per-record exec, vectorization, compilation

# Key Concepts: Relational Algebra

$\cap$ ,  $\cup$ ,  $-$ ,  $\times$ ,  $\sigma$ ,  $\Pi$ ,  $\bowtie$ ,  $G$

Algebraic rules involving these

# Key Concepts: Optimization

**Rule-based:** systematically replace some expressions with other expressions

**Cost-based:** propose several execution plans and pick best based on a **cost model**

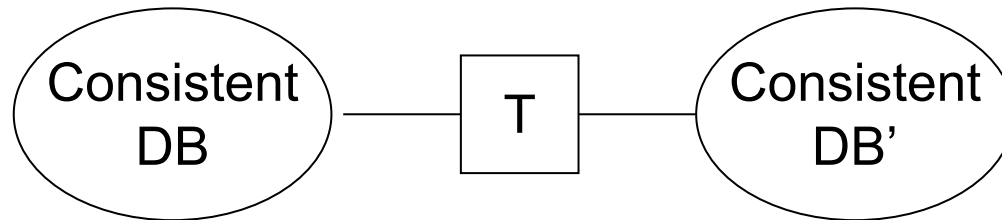
**Adaptive:** update execution plan at runtime

**Data statistics:** can be computed or estimated cheaply to guide decisions

# Key Concepts: Correctness

**Consistency constraints:** generic way to define correctness with Boolean predicates

**Transaction:** collection of actions that preserve consistency



**Transaction API:** commit, abort, etc

# **Key Concepts: Recovery**

**Failure models**

**Undo, redo, and undo/redo logging**

**Recovery rules** for various algorithms  
(including handling crashes during recovery)

**Checkpointing** and its effect on recovery

**External actions** → idempotence, 2PC

# Key Concepts: Concurrency

**Isolation levels**, especially **serializability**

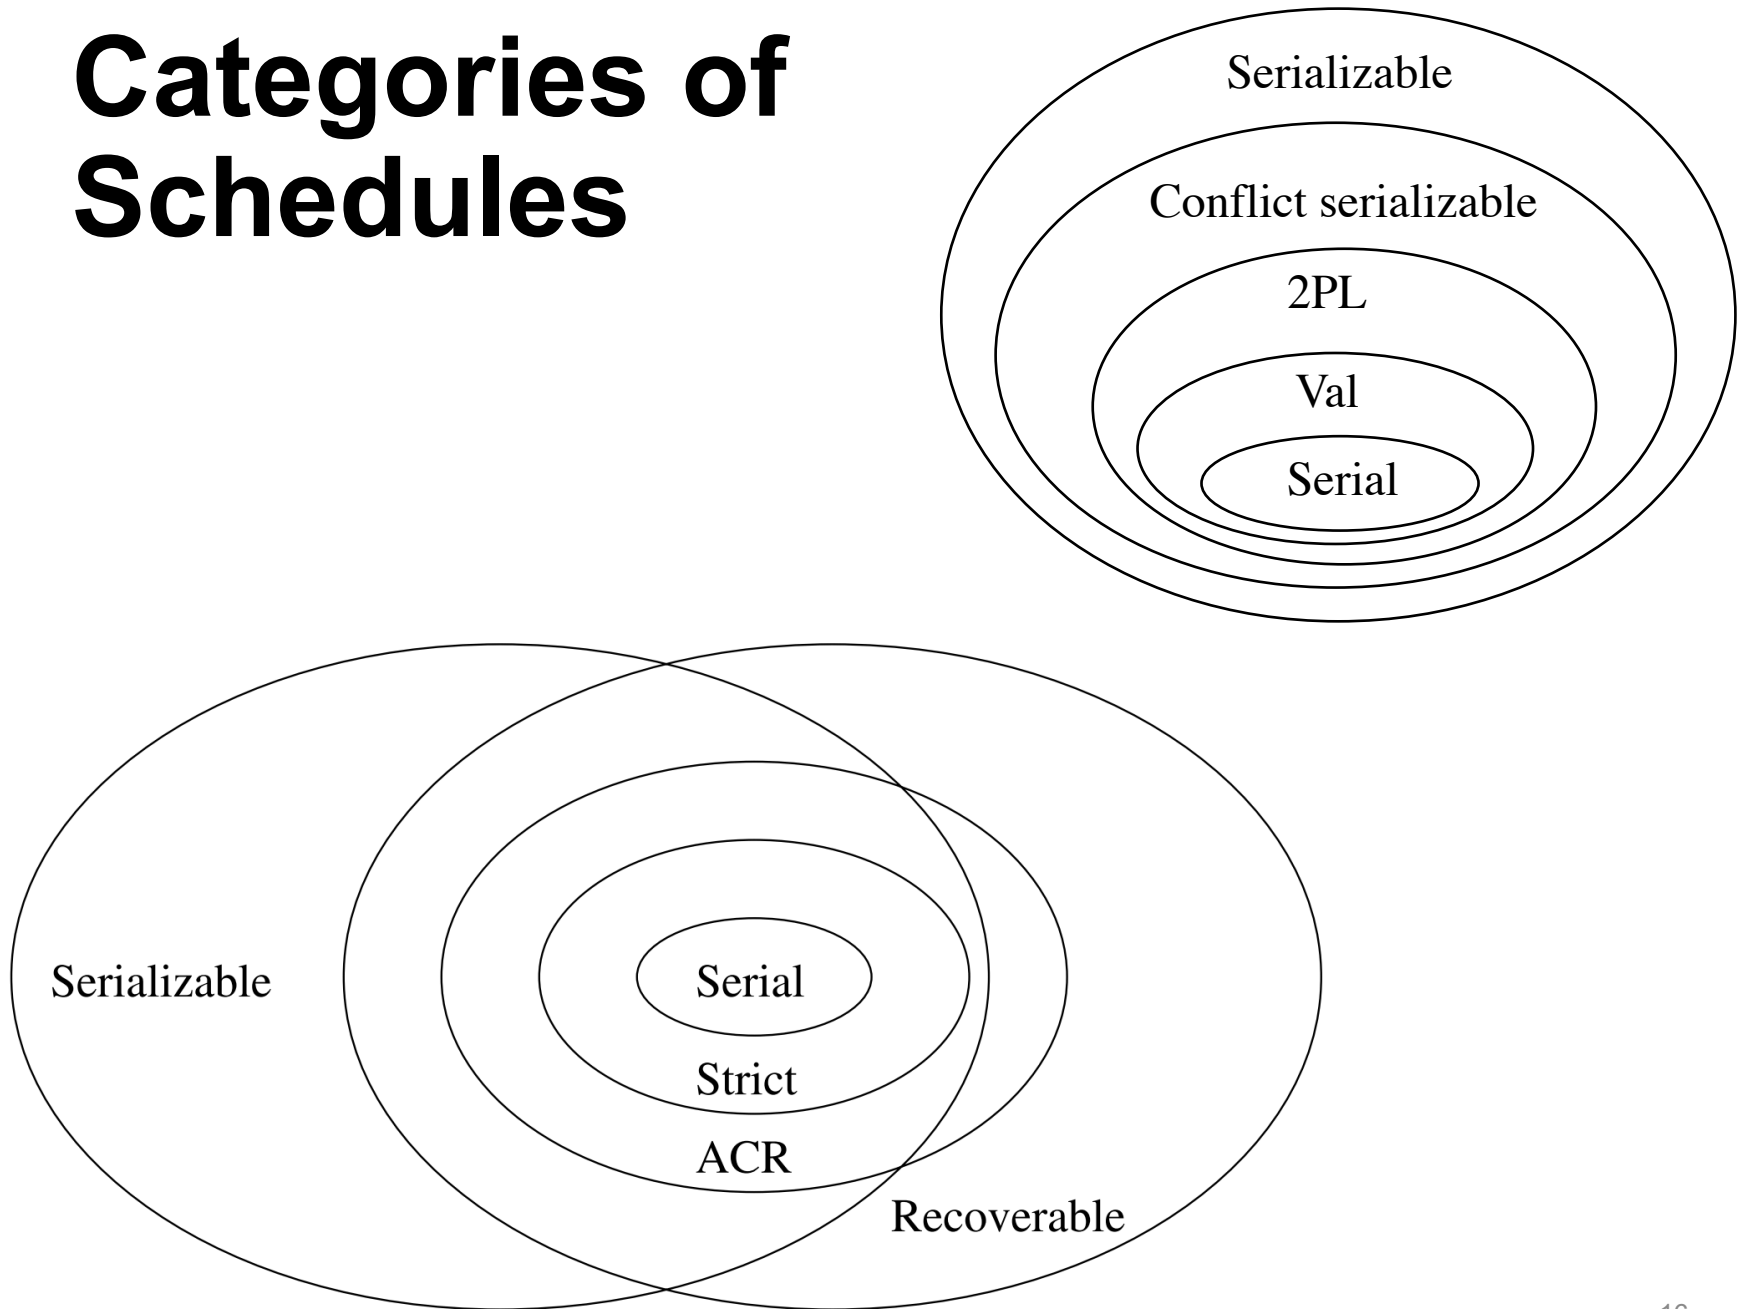
» Testing for serializability: conflict serializability, precedence graphs

**Locking:** lock modes, hierarchical locks, and lock schedules (well formed, legal, 2PL)

**Optimistic validation:** rules and pros+cons

**Recoverable, ACR & strict** schedules

# Categories of Schedules



# Key Concepts: Distributed

**Partitioning and replication**

**Consensus:** nodes eventually agree on one value despite up to  $F$  failures

**2-Phase commit:** parties all agree to commit unless one aborts (no permanent failures)

**Parallel queries:** comm cost, load balance, faults

**BASE** and relaxing consistency



# Key Concepts: Security and Data Privacy

## Threat models

**Security goals:** authentication, authorization, auditing, confidentiality, integrity etc

**Differential privacy:** definitions, computing sensitivity & stability

# Putting These Concepts Together

How can you integrate these different concepts into a coherent system design?

How to change system to meet various goals (performance, concurrency, security, etc)?

# **Send Us Your Feedback!**

We want to keep improving the course and tuning the content, so write a course eval